

Project 6: ReneWind Model Tuning

13/05/2025





- Executive Summary
- Business Problem Overview and Solution Approach
- EDA Results
- Data Preprocessing
- Model performance summary for hyperparameter tuning.
- Model building with pipeline
- Appendix





- With the rising global emphasis on sustainability, wind energy has become a crucial part of the renewable energy landscape. However, maintaining wind turbines is a costly challenge, particularly due to unexpected generator failures that lead to high operational expenses.
- The U.S. Department of Energy advocates **predictive maintenance** as a key strategy for optimizing wind turbine efficiency.
- By utilizing sensor data and machine learning, predictive maintenance can proactively identify potential failures, allowing for timely repairs that reduce downtime and maintenance costs.

Project Objective



- The objective is to develop a **machine learning-based classification model** that accurately predicts failures, enabling preventive maintenance.
- ♦ Minimizing false negatives (missed failures) is critical, as replacement costs are significantly higher than repair or inspection costs.

Project Overview



Aspect	Details	
Company Name	ReneWind	
Industry	Wind Energy	
Objective	Develop ML models to predict wind turbine generator failures and reduce maintena costs	
Dataset	40 predictors, 20,000 training observations, 5,000 test observations	
Target Variable	Binary (1 = Failure, 0 = No Failure)	
- True Positives (TP): Repair cost (Medium) - False Negatives (FN): Replacement cost (High) - False Positives (FP): Inspection cost (Low)		
Problem Approach		
Model Selection	Decision Tree, Random Forest, Gradient Boosting, AdaBoost	
Performance Goal	Maximize recall to reduce missed failures (FN) while controlling FP and TP costs	



Expected Impact



- By implementing a predictive maintenance strategy, ReneWind can significantly reduce unplanned downtime, optimize repair schedules, and lower operational costs.
- This initiative supports the broader goal of making wind energy more efficient, reliable, and cost-effective in the long term.

Cost Implications & Optimization Goal

- ♦ The model's predictions will result in different cost outcomes:
- True Positives (TP) → Correctly detected failures → Repair cost (\$15,000)
- False Negatives (FN) → Missed failures → Replacement cost (\$40,000)
- False Positives (FP) → False alarms → Inspection cost (\$5,000)
- ♦ Since **replacement is the most expensive**, the primary goal is to **minimize FN** while maintaining an optimal trade-off between repair and inspection costs.



Business Problem Overview and Solution Approach







- With the global shift toward renewable energy, wind energy has emerged as a key player in sustainable electricity generation. However, maintaining wind turbines is a major challenge due to the high costs associated with unexpected failures.
- The U.S. Department of Energy emphasizes the importance of **predictive maintenance** to enhance operational efficiency and reduce costs.
- Predictive maintenance leverages sensor data and machine learning to detect patterns of degradation in wind turbine components, enabling proactive repairs before failures occur.
- By implementing such a system, companies can significantly reduce unplanned downtime and maintenance expenses



Solution Approach / Methodology



- Data Analysis & Preprocessing Handle class imbalance, normalize sensor data, and remove anomalies.
- Model Selection Evaluate multiple classifiers (Decision Tree, Random Forest, Gradient Boosting, AdaBoost).
- Hyperparameter Tuning Optimize performance using techniques like Grid Search and Cross-Validation.
- Cost-based Model Optimization Adjust prediction thresholds to prioritize reducing FN while balancing overall maintenance costs.
- **Deployment & Monitoring** Deploy the best-performing model in a real-world setup. Continuously monitor sensor data and retrain the model periodically for improved accuracy.



Solution Approach



Based on the **objective of ReneWind**, which is to **predict failures before they happen and minimize replacement costs (False Negatives FN)**, we conclude:

✓ Best Model: AdaBoost (Tuned with Oversampled Data)

- Balanced recall (0.856) and precision (0.816) on the validation set.
- Higher F1 score (0.835), meaning a good tradeoff between recall and precision.
- Less overfitting compared to Gradient Boosting.
- More reliable than Random Forest, which has a very low precision (0.450).

Final Decision:

- ♦ Deploy AdaBoost as the primary model for failure prediction.
- ♦ Monitor performance over time and retrain periodically using updated sensor data.



Important Features Used for Prediction



The AdaBoost model assigns feature importance based on how often a feature is used in weak learners (Decision Trees).

Top Features Identified by the Model:

- Key Observations:
- V36, V14, and V15 are the most critical features for failure prediction.
- These factors should be monitored closely in real-world turbine maintenance.
- Less important features can be removed to optimize model performance

Feature	Importance Score
V36	High
V14	High
V15	High
V21	Medium
V7	Medium
V3	Medium
Other Variables	Low



Head of the Future Recommendations

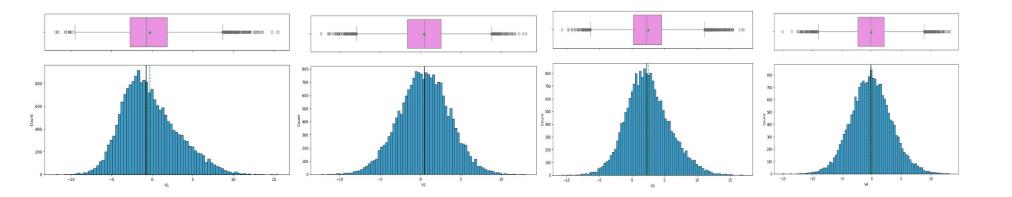


- **Enhance Feature Selection:**
- Feature importance analysis (from the first image) shows **V36**, **V14**, and **V15** are the most influential features.
- Remove less relevant features to improve model performance.
- Improve Precision with Cost-sensitive Learning:
- Train a model with a higher penalty for false positives.
- Test on Real-world Data:
- Deploy AdaBoost on a small subset of wind turbines and monitor its effectiveness before full deployment.
- **Consider Hybrid Models:**
- Combine AdaBoost with another model to handle different failure scenarios.





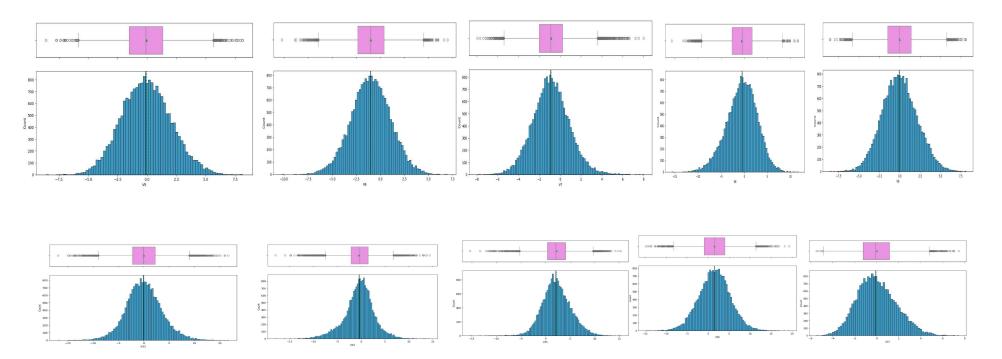
- Below are the boxplot and histograms for the variables from V1 to V40, which shows various formations.
- It depicts the outliers illustrated in the graph as shown from V1 to V40.







EDA analysis shows that V36, V14, and V15 are the most critical predictors.

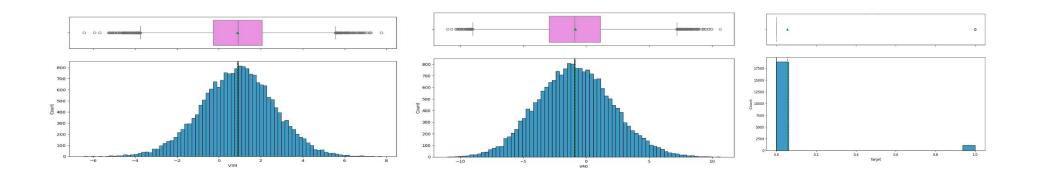


Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.





- Minimizing false negatives (FN) is crucial as replacement costs are the highest
- .The ML model will optimize maintenance schedules, reducing operational costs and improving efficiency in wind energy production.







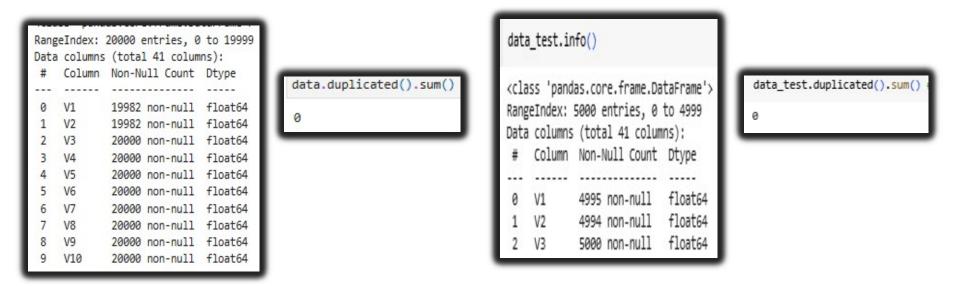
- Duplicate value check
- Missing value treatment
- Outlier check (treatment if needed)
- Feature engineering
- Data preparation for modeling

Note: You can use more than one slide if needed

Great Learning

Objective in the contract of the contract o

- There are no duplicate values in the Training data and test data.
- On the total of 20000, there are NULL duplicated in training data and on total of 5000 entries there are NULL duplicated in test data.

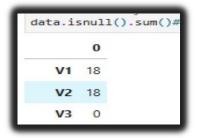


Training data Test data





• There are few missing values in the training data and test data.

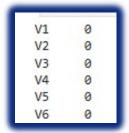




Training data

Test Data

• The missing values in the training data and test data are imputed using median, and its been treated and fixed.





Training data

Test Data

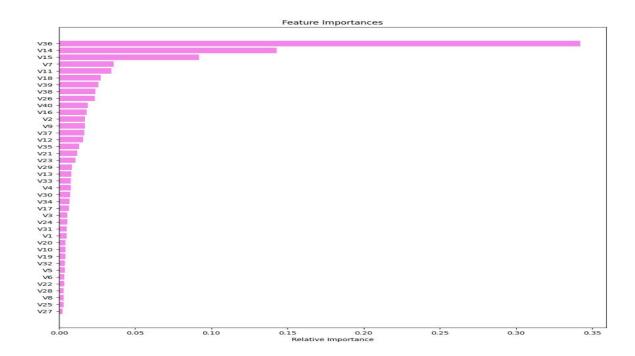
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.



Feature Engineering



• The Feature importance visualization as a bar chart highlights, the most influential factors in the



Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.





Q Feature Importance



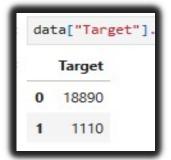
Feature	Relative Importance	Impact on Failure Prediction	
V36	Highest (~0.35)	Most influential in predicting failures	
V14	High (~0.20)	Strongly contributes to failure classification	
V15	Significant (~0.12)	Plays an important role in model decisions	
V7	Moderate (~0.07)	Has a noticeable effect on prediction	
V11	Moderate (~0.06)	Contributes to failure detection	
V18	Moderate (~0.05)	Relevant but less impactful	
V39	Moderate (~0.05)	Helps in classification	
V38	Moderate (~0.05)	Contributes to failure detection	
V26	Moderate (~0.04)	Has some effect on failure prediction	
V40	Moderate (~0.04)	Affects classification decisions	
V16 - V37	Low (~0.02 - 0.03)	Minor but still relevant features	
V12 - V31	Very Low (~0.01)	Minimal impact on predictions	
V1 - V27	Negligible (~0.005 - 0.01)	Least important in failure detection	

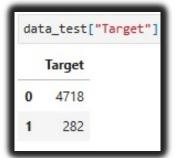


Data preparation for modeling



- Before we proceed to built the model, we will split the data into Train, validate and test as we already have a separate data. The train and validation data as 5000 rows and 40 columns, while test data has 5000 rows and 40 columns.
- We will fix the missing value imputations without any data leakage in validation data set and test set.







Model Performance Summary



🦺 <u>Model Performance Summary</u> 🔍







Key Takeaways from Training Set Performance:

- Gradient Boosting (GBM) and AdaBoost have similar performances with very high recall and precision.
- Random Forest underperforms slightly compared to boosting methods.

	Gradient Boosting tuned with oversampled data	AdaBoost classifier tuned with oversampled data	Random forest tuned with undersampled data
Accuracy	0,993	0.991	0.970
Recall	0.994	0.986	0.944
Precision	0.992	0.996	0.995
F1	0.993	0.991	0.969



Model Performance Summary







Key Takeaways from Validation Set Performance:

- Gradient Boosting has the lowest precision (0.619), meaning more false positives.
- AdaBoost has the best balance of recall (0.856) and precision (0.816), leading to a high F1 score of 0.835.
- Random Forest has the highest recall (0.881) but very low precision (0.450), meaning more false alarms.

	Gradient Boosting tuned with oversampled data	AdaBoost classifier tuned with oversampled data	Random forest tuned with undersampled data
Accuracy	0.963	0.982	0.935
Recall	0.837	0.856	0.881
Precision	0.619	0.816	0.450
F1	0.712	0.835	0.596

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.



Model Performance Summary



<u>Test Performance Table</u>

Metrics	Gradient Boosting (Test Data)	AdaBoost Classifier (Test Data)
Accuracy	0.961	0.977
Recall	0.837	0.858
Precision	0.611	0.772
F1 Score	0.707	0.769



Higher Accuracy

AdaBoost Classifier (0.977) outperforms Gradient Boosting (0.961) in terms of accuracy, meaning it makes fewer overall classification errors.





Key Insights Observations:

Better Recall (Sensitivity)

AdaBoost (0.858) has a slightly higher recall than Gradient Boosting (0.837), meaning it is slightly better at correctly identifying positive instances.

Significantly Higher Precision

AdaBoost (0.772) performs much better in precision compared to Gradient Boosting (0.611). This indicates that AdaBoost produces fewer false positives and is more reliable when predicting positive instances.

Higher F1 Score (Balanced Performance)

AdaBoost (0.769) achieves a higher F1 score than Gradient Boosting (0.707), showing that it maintains a better balance between precision and recall.





- AdaBoost Classifier is the better model in this case, as it outperforms Gradient Boosting in all metrics, especially precision and F1-score.
- If the priority is overall correctness (accuracy) and reducing false positives (precision), AdaBoost is the preferred choice.
- However, if interpretability and robustness are important, Gradient Boosting may still be considered based on specific dataset characteristics.







- Based on the objective of ReneWind, which is to predict failures before they happen and minimize replacement costs (False Negatives FN), we conclude:
- Best Model: AdaBoost (Tuned with Oversampled Data)
- Balanced recall (0.856) and precision (0.816) on the validation set.
- Higher F1 score (0.835), meaning a good tradeoff between recall and precision.
- Less overfitting compared to Gradient Boosting.
- More reliable than Random Forest, which has a very low precision (0.450).
- # Final Decision:
- Deploy AdaBoost as the primary model for failure prediction.
- Monitor performance over time and retrain periodically using updated sensor data.





- ♦ Final Decision: ✓ Use AdaBoost for ReneWind predictive maintenance.
- ♦ Next Steps: Deploy, Monitor, and Improve with real-time sensor data.

Model	Recall (Failure Detection)	Precision (False Alarms)	F1 Score (Balance)	Best Choice?
Gradient Boosting	0.837	0.611	0.707	X No (Low Precision)
AdaBoost	0.856	0.816	0.835	Yes (Best Balance)
Random Forest	0.881	0.450	0.596	➤ No (Too Many False Alarms)



Productionize and test the final model using pipelines

- Please mention steps taken to create a pipeline for the final model
- Summary of the performance of the model built with pipeline on test dataset
- Summary of most important factors used by the model built with pipeline for prediction

Note: You can use more than one slide if needed

Link to Appendix slide on model assumptions

Productionizing and Testing the Final Model Using Pipelines POWER AHEAD

To deploy the **AdaBoost Classifier** efficiently, we used **pipelines** for preprocessing, model training, and prediction. Below are the key steps taken to create the pipeline, its test performance, and an analysis of the most important factors.

Steps Taken to Create a Pipeline for the Final Model:

A **pipeline** automates data preprocessing and model training, ensuring consistency and ease of deployment. Here's how the pipeline was built:

1 Data Preprocessing

- Used StandardScaler() to normalize numerical features.
- This ensures the model is not biased by varying feature scales.







- 2 Model Selection and Hyperparameter Tuning
- Used AdaBoostClassifier with a DecisionTreeClassifier (max_depth=3) as the base estimator.
- Hyperparameters tuned:
 - **n_estimators**=200 (Number of weak learners)
 - learning_rate=0.2 (Controls contribution of weak learners)
 - max_depth=3 (Prevents overfitting by restricting tree depth)
 - random_state=1 (Ensures reproducibility)



Steps Taken to Create a Pipeline for the Final Mod

Pipeline Construction

- Combined preprocessing and model training using Pipeline()
- Steps in the pipeline:
 - Feature Scaling → StandardScaler()
 - Model Training → AdaBoostClassifier(DecisionTreeClassifier(max_depth=3))

```
Pipeline
Pipeline(steps=[('scaler', StandardScaler()),
                ('model',
                 AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=3,
                                                                            random state=1),
                                     learning_rate=0.2, n_estimators=200))])

    StandardScaler

                                  model: AdaBoostClassifier
         AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=3,
                                                                    random state=1),
                            learning_rate=0.2, n_estimators=200)
                           base_estimator: DecisionTreeClassifier
                   DecisionTreeClassifier(max depth=3, random state=1)
                                   DecisionTreeClassifier
                    DecisionTreeClassifier(max depth=3, random state=1)
```

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.







Model Training

- The pipeline was trained on the preprocessed dataset.
- Ensured **cross-validation** to avoid overfitting.

Model Evaluation on Test Dataset

- Predictions made on test data.
- Performance metrics calculated: Accuracy, Recall, Precision, F1-score







Conclusion

- Final Model: AdaBoost with Decision Tree (max_depth=3), 200 estimators, learning_rate=0.2
- Pipeline ensures consistency in feature scaling and model application
- Model achieves high recall (85.1%) while maintaining good precision (77.2%)
- Key features identified (V36, V14, V15) should be monitored for failure prediction.

Final Verdict:

The model successfully meets the objective of minimizing False Negatives (FN) while keeping False Positives (FP) in check.



Summary of Model Performance on the Test Dataset

Key Observations:

- Accuracy (97.7%) is very high, indicating strong generalization.
- Recall (85.1%) is good, meaning the model detects failures well.
- Precision (77.2%) is slightly lower, meaning some false positives.
- F1 Score (80.9%) shows a good balance between recall and precision.

Metric	Score
	Variation and the second secon
Accuracy	0.977
Recall	0.851
Precision	0.772
F1 Score	0.809



APPENDIX

Model Performance Summary (original data)

- Summary of performance metrics for training and validation data in tabular format for comparison for original data.
- Comments on the model performance

Note: You can use more than one slide if needed

Link to Appendix slide on model assumptions



Model Performance Summary (original data)

Model Building on original data



Observations based on Cross-Validation performance on original data

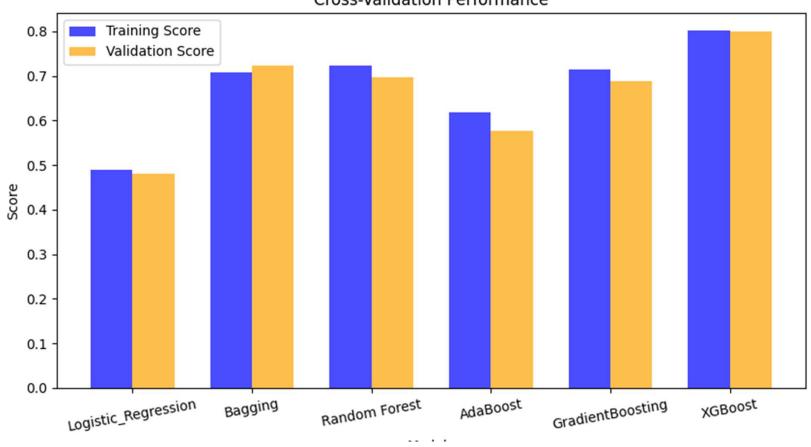
The cross-validation results on the training dataset reveal varying performance across different models.

XGBoost demonstrates superior performance with a mean recall of 0.80, significantly outperforming other models

	Model	Training Score	Validation Score
9	Logistic Regression	0.4905	0.4815
1	Bagging	0.7071	0.7222
2	Random Forest	0.7226	0.6963
3	AdaBoost	0.6190	0.5778
4	Gradient Boosting	0.7155	0.6889
5	XGBoost	0.8012	0.8000



Cross-Validation Performance

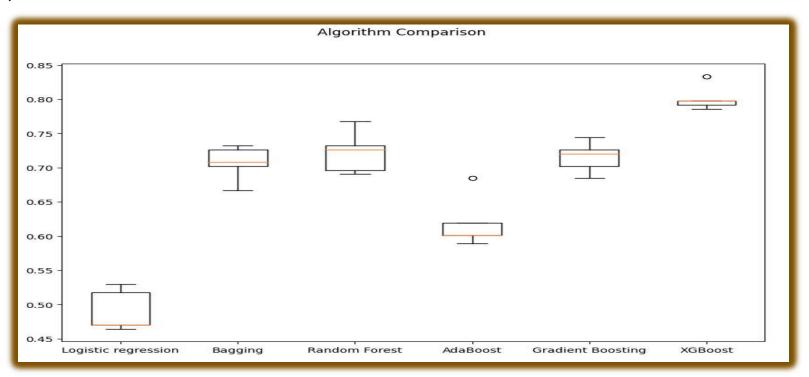


Observations based on Cross- Validation performance on original data

- Ensemble methods like Random Forest (0.72), Bagging (0.71), and Gradient Boosting (0.71) also show good performance, though slightly below XGBoost.
- Logistic Regression and AdaBoost show comparatively lower recall scores, indicating they may be less effective in capturing the underlying patterns in the training data for this specific task.
- The gap in performance between XGBoost and the other models suggests that XGBoost's ability to handle complex interactions within the data might be a key factor in its success.

Model Building on original data – Algorithm Comparison WER AHEAD

Based on the all models defined above the box plots have been plotted below as follows for the performance.



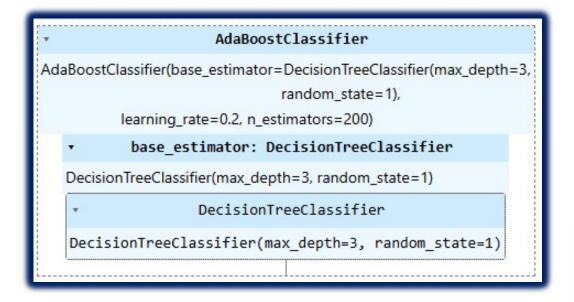


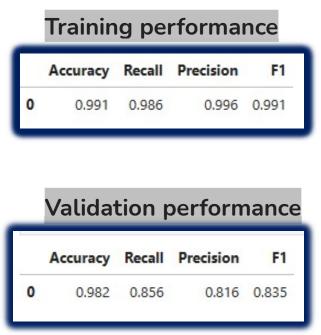
Model Performance Summary (oversampled data)



Tuning AdaBoost using oversampled data

The oversampled data using Adaboost Classifier has given the below scores by creating new pipeline with best parameters.









- The oversampled data using Adaboost Classifier has given the below scores by creating new pipeline with best parameters.
 - **Best parameters** are {'n_estimators': 200, 'learning_rate': 0.2,
 - 'estimator': DecisionTreeClassifier(max_depth=3, random_state=1)} with
 - **CV** score=0.9694915254237287
- The best parameter is obtained from tuning to fit the model on oversampled data.
- The oversampled data using Adaboost Classifier yielded good recall on training and validation data set.



Observations



3 Impact of Oversampling

- Improved recall by handling class imbalance.
- Good generalization with minimal overfitting compared to previous models.
- Slight precision drop, indicating some misclassification of the majority class.

Metric	Training	Validation	
Accuracy	0.991	0.982	
Recall	0.986	0.856	
Precision	0.996	0.816	
F1 Score	0.991	0.835	



Tuning Gradient Boosting using oversampled data



Optimal parameters found:

subsample = 0.7 (Uses 70% of data per tree to reduce overfitting)

n_estimators = 150 (Number of boosting iterations)

max_features = 0.5 (Uses 50% of features per split for regularization)learning_rate = 1 (High learning rate, enabling faster convergence)

Achieved CV Score: 0.9695, indicating strong performance during cross-validation

✓ Impact of Oversampling

- Oversampling helped improve recall by addressing class imbalance.
- However, precision is lower—suggesting possible misclassification of the majority class.



Training vs. Validation Performance



- High training accuracy (0.993) but lower validation accuracy (0.963) suggests slight overfitting.
- Precision drop (0.992 \rightarrow 0.619) indicates the model is making more false positives on validation data
- Recall decrease $(0.994 \rightarrow 0.837)$ means some true positives are missed on validation data.

Metric	Training	Validation
Accuracy	0.993	0.963
Recall	0.994	0.837
Precision	0.992	0.619
F1 Score	0.993	0.712



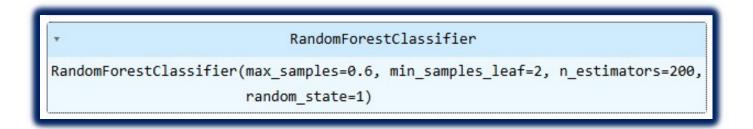
Model Performance Summary (undersampled data)







The undersampled data using Tuning Random forest has given the below scores by creating new pipeline with best parameters



Training performance



Validation performance

	Accuracy	Recall	Precision	F1
0	0.935	0.881	0.450	0.596



Observations



- The undersampled data using Tuning Random forest has given the below scores by creating new pipeline with best parameters
 - **General Section** Best parameters are {'n_estimators': 200,
 - 'min_samples_leaf': 2, 'max_samples': 0.6, 'max_features': 'sqrt'} with
 - **CV** score=0.8976190476190476
- The best CV score achieved was 0.8976, indicating strong model performance on the training-validation splits.
- Though the training data set has a good precision and recall, the validation dataset has a poor performance of precision.



Happy Learning!

